

CS747:Programming Assignment 1 Report

Debabrata Biswal
Roll no : 203050024
September 25, 2020

1 Introduction

The report contains the implementation details of the algorithms that I have coded for this assignment including the assumptions, some observations. P.S : All notations are according to the class conventions. So I am not defining them separately.

2 Epsilon Greedy

The epsilon was taken to be 0.02 as mentioned in the assignment. I am directly starting the algorithm without pulling each arm for some fixed times, as I think that's not really necessary for this algorithm. The regret is linear as it should be.

3 UCB

This was also a straightforward algorithm to implement. Although here I am pulling each arm once just so that for every arm u_t^a is non-zero as that will otherwise cause problem in calculation of upper confidence. For this algo I am setting upper confidence bound for arm a as

$$ucb_a^t = \hat{p}_a^t + \sqrt{\frac{2\ln(t)}{u_a^t}}$$

The regret for this algorithm was increasing logarithmically as should be the case.

4 KL-UCB

I really enjoyed implementing this algorithm, especially the part of optimizing this algorithm to decrease the time from 4 minute to just 45/55 seconds for instance 3 and horizon 102400. (All thanks to the vectorization). Firstly I am pulling each arm once before starting again just for convenience of calculations. Now since I pulled each arm once then empirical mean of each arm will now be either 0 or 1. But to avoid some obvious pitfalls (like $\log(0)$ is $-\infty$ and $KL(p_a^t, 1)$ is ∞) I have added very small amount ($1e-12$) to the arm having

empirical mean 0 and subtracted 1e-12 from arms having empirical mean 1. This was performing better than UCB algorithm as it is known to have smaller constant coefficient and tighter upper confidence bound. I used $C = 3$ for this algo although having C less than 3 is performing better practically.

5 Thompson Sampling

Again This was also a fairly Simple and Straight Forward algorithm. We set a prior for mean of the each arm which is a beta distribution. Initially parameters of the beta are set to be $\alpha = 1$ and $\beta = 1$. Then we update our posterior distribution according to the reward. This was found to be performing better than KL-UCB for every instance.

6 Thompson Sampling With Hint

I learnt so many things solving this problem, especially I got clarity of how Thompson sampling works. Now the problem was to come up with a algorithm which performs better than Thompson sampling if you know the means of the arms beforehand but just the permuted means.

I first tried some algorithms based on the ideas of Median Elimination and all, But It seemed like I am not using much of the information that is being provided to me. I was just using the knowledge of Maximum mean but not the whole permuted means that is being given to me.

So after some thought I decided do just like Thompson sampling , But instead of using a beta prior I will use a Discrete Prior. For each arm a there will be a discrete distribution that will have $P(\text{mean} = \theta | \text{arm} = a)$ for each θ from the array of permuted means. Now according to the reward I get from a arm I will update the distribution for that arm. As i said It is exactly equal to the Thompson sampling but with a discrete prior instead of a beta prior. The update rule is straight-forward.

$$p(\text{mean} = \theta | \text{reward}) = \frac{p(\text{reward} | \text{mean} = \theta) p(\text{mean} = \theta)}{\sum_j p(\text{reward} | \text{mean} = \theta_j) p(\text{mean} = \theta_j)}$$

Now this algorithm worked as expected, But due to discrete values the sampling dis not work as expected, So this algorithm still lacked a little behind Thompson sampling for some instances. Now being out of ideas , I thought

to again use the information we are provided. So when after some iteration the distribution of the optimal arm saturates to the point that $p(\text{mean} = \text{highest}\theta | \text{arm} = a)$ is above 0.99 I choose that arm and starting from that point I only pull the chosen arm. Now as expected This modification faired way better.

7 Just a observation

For instance 3 When I compared UCB with Epsilon-Greedy I noticed that even till Horizon 102400 Epsilon Greedy was performing better than UCB. Although if you look at the slopes you can clearly see that regret for epsilon-greedy was increasing much more rapidly. So just to confirm it I checked for horizon 500000 over 50 random seeds and averaged the regret and found the following.

The mean Regret for epsilon greedy was - 4814.32

The mean regret for epsilon greedy was - 2393.04

As expected UCB surpassed epsilon greedy.




Figure 1: Comparison of different algorithms for instance 1

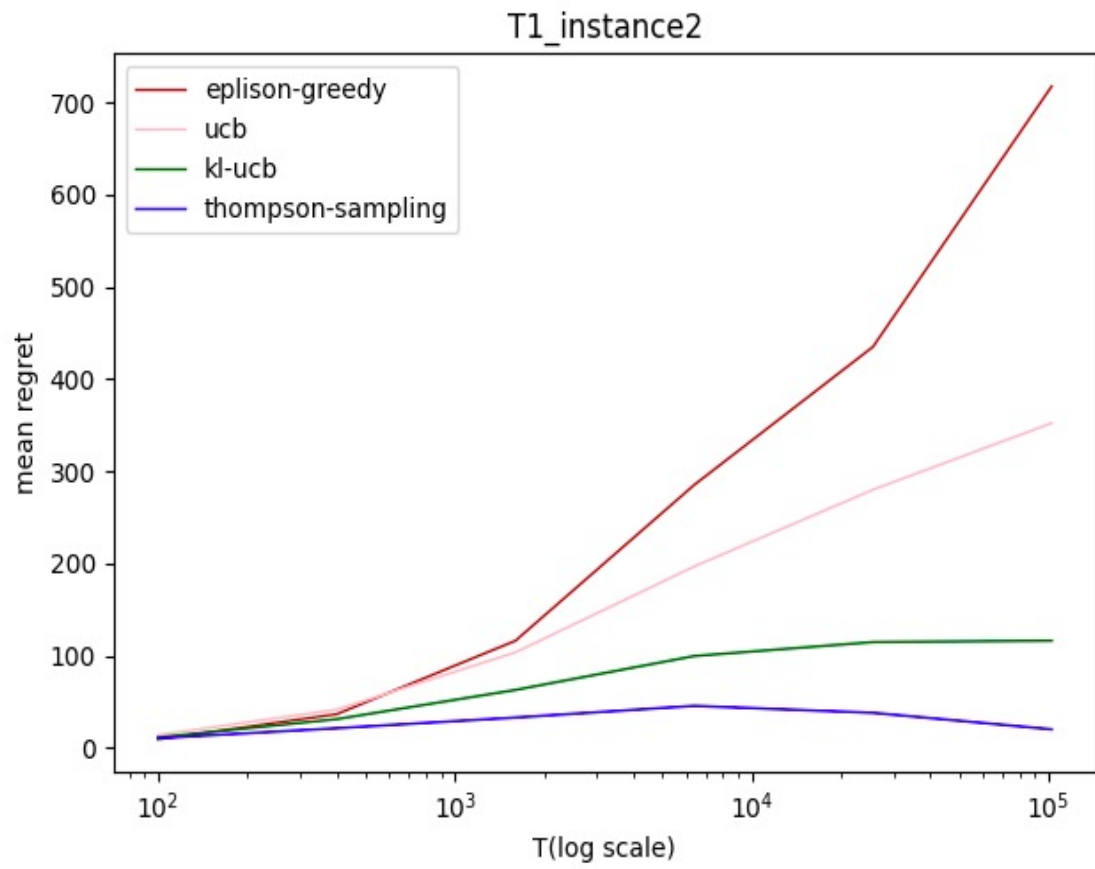


Figure 2: Comparison of different algorithms for instance 2

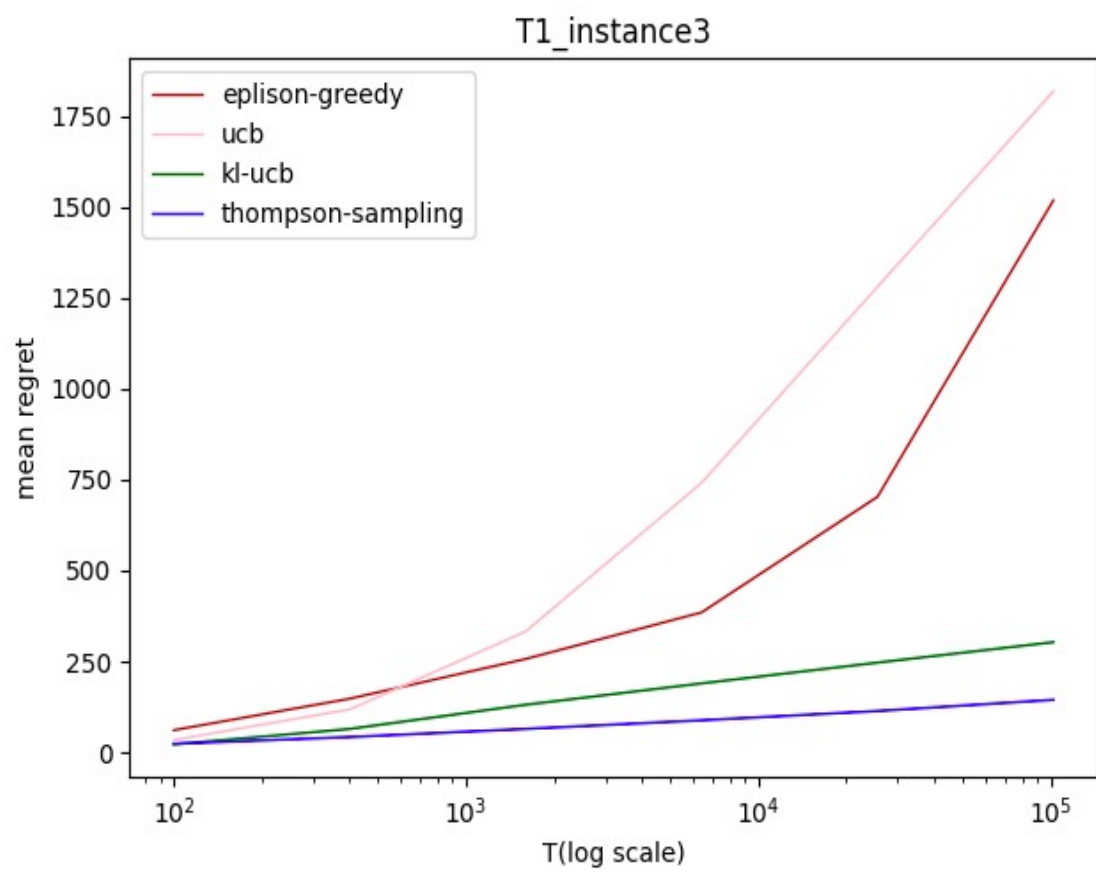


Figure 3: Comparison of different algorithms for instance 3

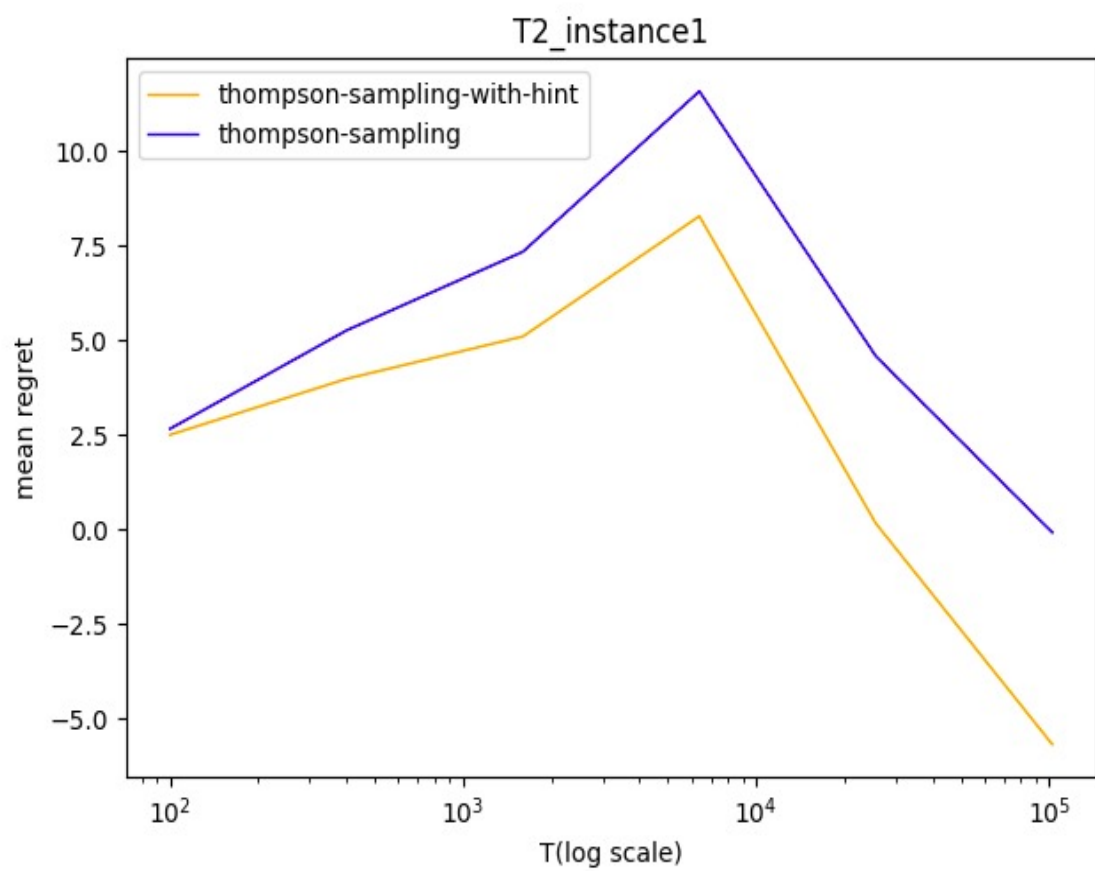


Figure 4: Comparison of Thompson Sampling and Thompson Sampling with hint for instance 1

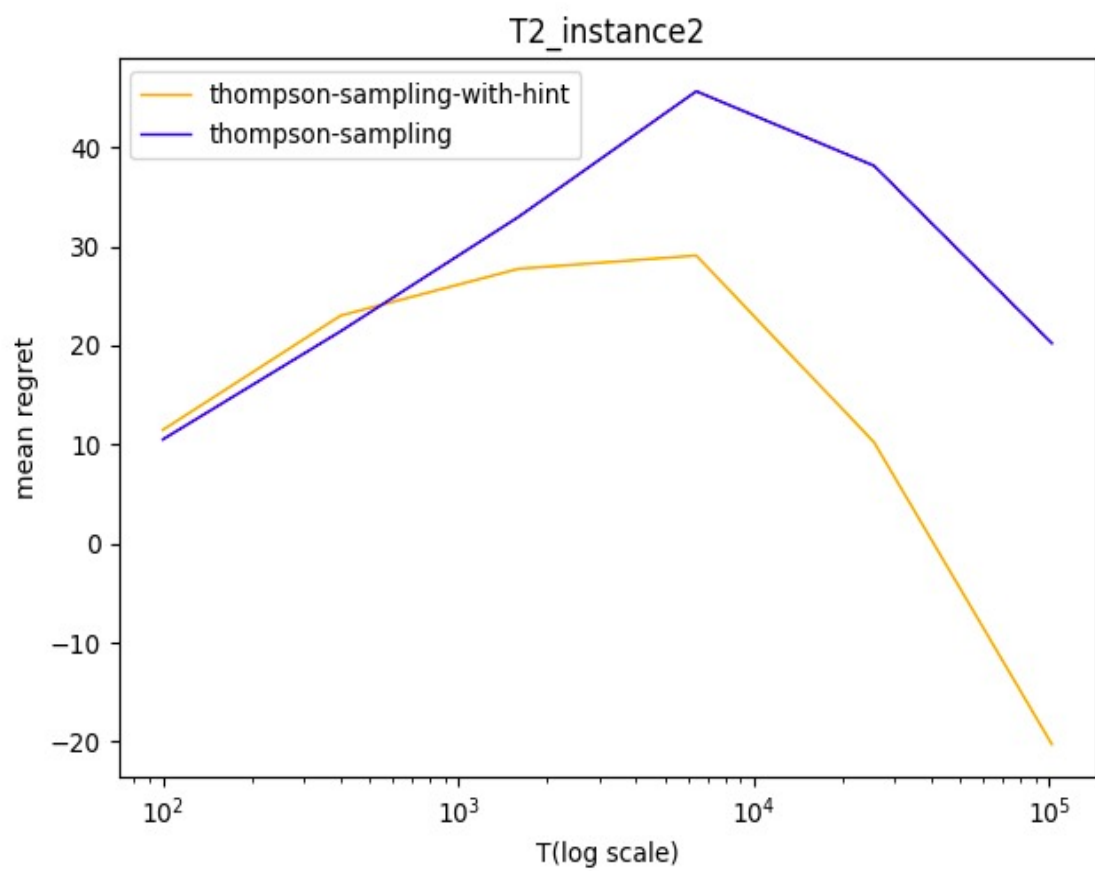


Figure 5: Comparison of Thompson Sampling and Thompson Sampling with hint for instance 2

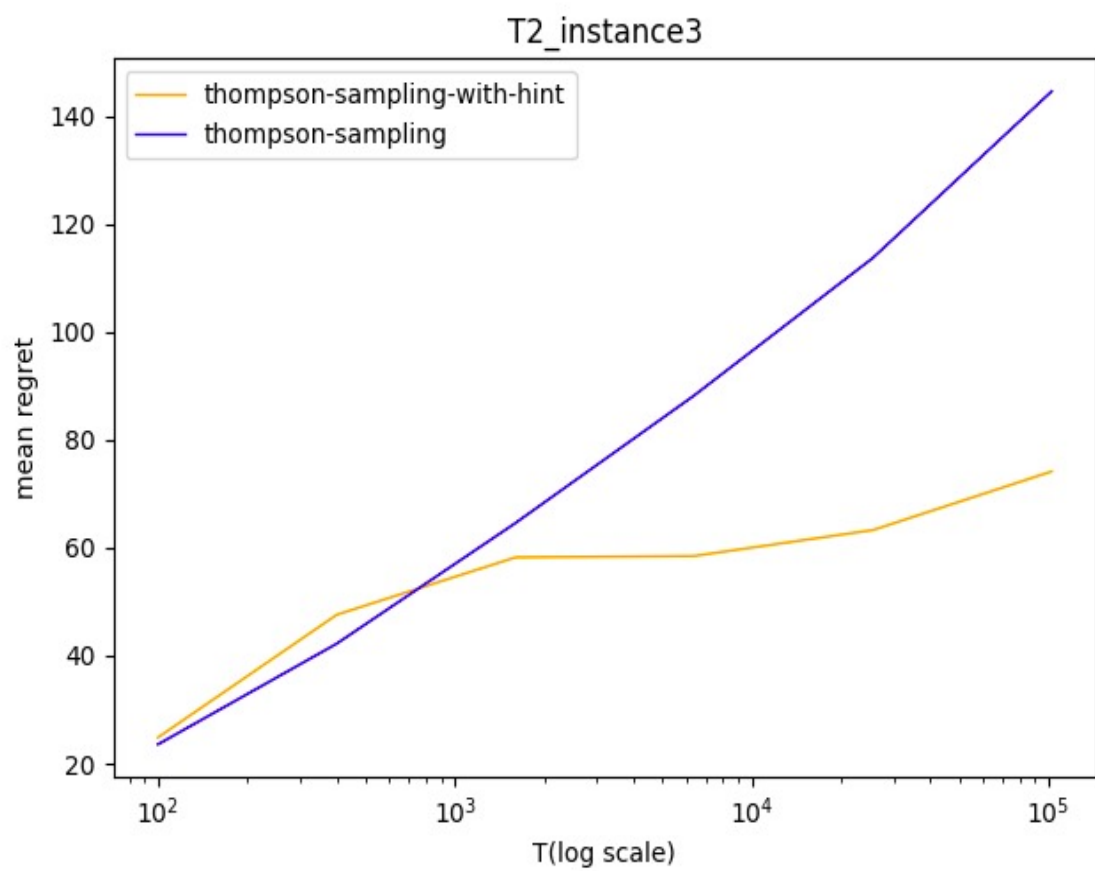


Figure 6: Comparison of Thompson Sampling and Thompson Sampling with hint for instance 3

8 T3

Here are ϵ values I found such that $\epsilon_1 < \epsilon_2 < \epsilon_3$ but regret corresponding to ϵ_2 is lowest.

for instance 1 : $\epsilon_1 = 0.001, \epsilon_2 = 0.005, \epsilon_3 = 0.1$

for instance 2: $\epsilon_1 = 0.0001, \epsilon_2 = 0.02, \epsilon_3 = 0.1$

for instance 3: $\epsilon_1 = 0.0001, \epsilon_2 = 0.02, \epsilon_3 = 0.1$